

Data Structures and Algorithms

Assignment

Semester 2, 2016

**Department of Computing
Curtin University**

1 Preamble

In practicals you have implemented and learned about a number of algorithms and ADTs and will be implementing more of these in the remaining practicals. In this assignment, you will be making use of this knowledge to implement a system that is somewhat more open ended. In this system, you will be deciding (for the most part) how to structure the program, how to use the abstract data types and your program will need to decide on the flow of objects in and out of them. Feel free to re-use the generic ADTs from your practicals. However, remember to self-cite; if you submit work that you have already submitted for a previous assessment (in this unit or any other) you have to specifically state this.

The start of this assignment specification document does look somewhat daunting and confusing. I recommend that you highlight and/or take notes. Jump to Section 4 for a breakdown of what your program is expected to do for each set of marks. Note that, as you can see from Section 4, you can get many of the marks for properly completing your practicals properly and then doing a small amount of work to apply that to this assignment.

1.1 Requirements for passing the unit

Please note: As specified in the unit outline, it is necessary to have attempted the assignment in order to pass the unit. As a guide, you should score at least 15% to be considered to have attempted this assignment. We have given you the exact mark breakdown in Section 4.2. Note that the marks indicated in this section represent maximums, achieved only if you completely satisfy the requirements of the relevant section.

Plagiarism is a serious offence. This assignment has many correct solutions so plagiarism will be easy for us to detect (and we will). For information about plagiarism, please refer to <http://academicintegrity.curtin.edu.au>.

In the case of doubt, you may be asked to explain your code and the reason for choices that you have made as part of coding to the unit coordinator. A failure to adequately display knowledge required to have produced the code will most likely result in being formally accused of cheating.

Finally, be sure to secure your code. If someone else gets access to your code for any reason (including because you left it on a lab machine, lost a USB drive

containing the code or put it on a public repository) you will be held partially responsible for any plagiarism that results.

1.2 Late Submission

I know for the vast majority of you, I don't need to tell you what is in this section. But every semester, someone always tries to see what they can get away with here. For the benefit of fairness to the vast majority of you who do the right thing, we need to make the rules surrounding late submission clear

As specified in the unit outline, you must submit the assignment on the due date. Acceptance of late submissions is not automatic and will require supporting documentation proving that the late submission was due to unexpected factors outside your control. See the unit outline for details as to the procedure for requesting that an assessment be accepted after the due date.

Note that external pre-scheduled commitments including, but not limited to, work, travel, scheduled medical, sporting, family or community engagements are not considered unexpected factors outside your control. If you know you have, or are likely to have, such engagements and that they may affect your ability to complete the assignment, you will be expected to have planned your work accordingly. This may mean that you need to start and/or complete your assignment early to make sure that you are able to hand it in on time.

Also note that IT related issues are almost never a valid excuse. These include computer crashes, hard disk corruption, viruses, losing computers/storage media, networks going down (even if it is Curtin's network - outages of the entire Curtin network of more than 24 hours may be considered depending on circumstances) or the like. As IT professionals in training, you are expected to have suitable backups and alternative ways of getting your assignment completed in the event that any IT problems are encountered. You are also expected to submit your assignment ahead of time to allow for unforeseen issues.

In the event that you submit your assignment late and are deemed to have a valid excuse, you will be penalised 10% (that is, 10% out of 100%, not out of what you would have received) per calendar day that you are late, up to a maximum of seven (7) calendar days. Any work submitted after this time will not be marked and you will automatically fail the unit. Note that if you are granted an extension you will be able to submit your work up to the extended time without penalty - this is different from submitting late.

Note that the requirements for passing this unit are applied after penalties. An assignment normally scoring 20% that is submitted one day late, even with a valid excuse, will score only 10% and thus not satisfy the requirements for passing this unit.

1.3 Clarifications and Amendments

This assignment specification may be clarified and/or amended at any time. Such clarifications and amendments will be announced in the lecture and on the unit's

Blackboard page (not necessarily at the same time and not necessarily in that order). These clarifications and amendments form part of the assignment specification and may include things that affect mark allocations or specific tasks. It is your responsibility to be aware of these, either by attending the lectures, watching the iLecture and/or monitoring the Blackboard page.

Note that there will be at least one minor amendment to these requirements. In a normal real-world project a large number of amendments and clarifications would be made, so marking at least one minor amendment is helpful in teaching you to adjust. The planned amendment(s) will not require a large amount of time assuming you have made sensible decisions.

2 The Problem

Your assignment will address the problem of computing the reliability of a network. In particular, you will compute the two-terminal reliability of a directed network. This is surprisingly hard to do; all existing solutions have an exponential worst-case complexity. Compare this with the algorithms that have been covered so far that generally have quadratic complexity at worst. For comparison, a quadratic with input 100 (in other words we compute $100^2=10,000$) isn't very large, but even a small exponential function with input 100 is huge (2^{100} has about 30 digits). In other word it's hard to compute the two-terminal reliability of even a network with only 100 devices. In this assignment you will be computing the reliability of networks with thousands of devices.

The first part of the problem is to read the network in from a provided file. These files are discussed in Appendix II. Once this is done, the data may need to be sorted as specified in Appendix I and the specifications for the Network class in Section 3.4.

The second part of the problem is to process the network, one vertex at a time, to compute the reliability value. The main algorithm in the OBDDH class will be provided to you, as will the details of HNode creation and modification which make up the most challenging parts of the algorithm. The details of both of these classes can be found in Section 3. The OBDDH class represents the main decision diagram used to compute the reliability and each node of that diagram is a HNode object.

Your job is to provide the support classes that will allow this to function. This simulates working in a team with an expert in the area who will supply the algorithm-specific code while you provide the production code.

As part of the teamwork simulation, you are required to write code for a number of classes and other code will be provided for you. This is specified in Section 3. The code provided will be released in parts, to simulate your team member's progress on their work. This means that you'll have to do coding on your sections without being able to test how your code fits in with the provided code. The only way to do this is for both parties to stick to specifications exactly and to carefully unit test their code.

3 Class Specifications

This section contains the specifications of the main classes required in this assignment. Some of the code for this will be given to you over time (this represents code created by other members of your programming team). All code given to you will adhere to the specifications in this section, meaning you have to be careful that your code does as well or you will not be able to fit the various pieces of code together. You may not change any part of the code given to you.

In order to implement the classes that are your responsibility, you may need to create other classes. Decisions such as this (and all main decisions) must be documented – see Section 4.1. You will be marked not just on having working code but also on the efficiency of the overall solution (part of which is making good decisions) and the quality of your code.

Each class lists the methods that will be provided and those that you are required to write. Provided methods will be made available over time and required methods are ones that will be utilized by the provided methods; you must code required methods to exactly match the specifications or they will not work with the provided code. You may also need to code other methods that are not listed in order to use the classes yourself. Your implementation of all classes listed below must be named exactly as specified.

Note that while the classes must implement the required methods, it is acceptable for these to be inherited from a parent class where appropriate. In other words the calls to those methods must function but the actual code can be elsewhere, at your discretion. When to use inheritance is an important part of OO programming, and you should identify this in your documentation (see Section 4.1).

You must write the classes yourself. You may not use existing Java classes for any class listed below.

3.1 FloatingPoint

This ADT is responsible for safe handling of floating point numbers. This class should have a static class variable *precision* that is of type double. This should be set to 0.00000001 by default.

Provided: None

Required:

- public FloatingPoint(double value) – constructor that sets the initial contents of the FloatingPoint to the double value
- public FloatingPoint(FloatingPoint num) – copy constructor
- public boolean equals(FloatingPoint num) – checks to see whether the two numbers are within *precision* of each other.
- public static void setPrecision(double value) – sets the precision to the new value.

- `public void add(FloatingPoint num)` - adds the given `FloatingPoint` number to the current one. So calling `a.add(b)` will result in `a` containing the sum of `a` and `b`.
- `public void add(int num)` - adds the given integer to the current `FloatingPoint` one. So calling `a.add(b)` will result in `a` containing the sum of `a` and `b`.
- `public void invert()` - inverts the floating point number using probability laws by subtracting it from 1. So calling `a.invert()` will result in `a` containing the result of $1.0 - a$.
- `public void multiply(FloatingPoint num)` - multiplies the given `FloatingPoint` number into the current one. So calling `a.multiply(b)` will result in the object `a` containing the product of the values previously stored in objects `a` and `b`. The value in object `b` should not be changed.
- `public String toString()` - returns the floating point number stored in the class as a `String`

3.2 Vertex

This ADT is responsible for storing the devices of the network as vertices in the graph. Vertices must store all of the information related to them from the network file (see Appendix II).

Devices must have one of the following device types: RV325, 2291K9, 1841

Provided: None

Required:

- `public FloatingPoint getRel()` - returns the reliability value stored in the vertex as a member of the `FloatingPoint` class.
- `public boolean isSource()` - true if the vertex is the source vertex (as defined in the file specifications in Appendix II)
- `public boolean isTarget()` - true if the vertex is the target vertex (as defined in the file specifications in Appendix II)
- `public String getName()` - returns the name of the device
- `public String getType()` - returns the type of the device
- `public String toString()` - convert the `Vertex` into a `String` for display.

3.3 Edge

This ADT is responsible for storing the communications of the network as edges in the graph. Edges must store all of the information related to them from the network file (see Appendix II).

Edges must have one of the following types: CAT4, CAT5, Wireless, Satellite

Provided: None

Required:

- `public FloatingPoint getRel ()` - returns the reliability value stored in the edge as a member of the `FloatingPoint` class.

- `public String getName()` – returns the name of the device
- `public String getType()` – returns the type of the device
- `public Vertex getFrom(), public Vertex getTo()` – get the endpoints of the edge.
- `public boolean isDirected()` – returns *true* if the edge is directed and *false* if it is undirected.
- `public String toString()` – convert the Edge to a String for display.

3.4 Network

This class is the ADT responsible for managing the network to be analysed. A network contains vertices and edges, which must be stored in the same order as read in from a sorted file (with suffix `.srt`) or otherwise they must be sorted before being used.

Provided: None

Required:

- `public Network()` – empty constructor.
- [`Public int getEdgeCount\(\)` – returns the number of edges in the network](#)
- `public Vertex nextVertex()` – returns the next vertex stored in the network. This is used to traverse the vertices one at a time, in the order that results from the sorting. The first time this method is called the first (source) vertex should be returned, and when no more vertices are available a null should be returned.
- `public Vertex[] adjacent(Vertex v)` – returns all vertices adjacent to the current vertex. A vertex *w* is adjacent to vertex *v* if an edge exists that has *v* as the *from* vertex and *w* as the *to* vertex.
- `public Edge[] adjacentE(Vertex v)` – as above, but returns all of the edges adjacent to *v* instead of the vertices. An edge is adjacent to *v* if it has *v* as an endpoint (either to or from). The size of the array should be correct (ie if there are three adjacent edges then the array returned has size 3).
- `public Network(String filename)` – a constructor that accepts the file name and reads the contents of the network in from a network file (see the specification in Appendix II). Throws appropriate IO exceptions. If the network file ends with the `.srt` suffix, it is assumed to be already sorted. If this is not the case (it should end in a `.nt` suffix) then it should be sorted using the `sort()` method below and then saved. The filename in this case is identical but the `.nt` suffix will be replaced by a `.srt` suffix. In other words, if you load the file “file.nt” the fact that it’s unsorted should be detected, the network should be sorted, and the sorted network should be written to the file “file.srt”. Sorting is described in Appendix I.
- `public void display()` – write the network to the screen in a sensible fashion.
- `private void sort()` – sort the network. Vertices must be sorted in increasing distance from the source vertex and edges must be sorted in increasing order for the lower of the two endpoints and then the upper endpoint. See Appendix I for details on vertex sorting.

3.5 HNode

This is a node of the hybrid decision diagram. It has a private class variable *FloatingPoint reliability* that stores the reliability associated with the network state represented by the node. It also has a private class variable *state* that stores the state the node represents.

Provided:

- public HNode() – the default constructor creates the root node of the OBDDH diagram
- public boolean isEqual(HNode node) – returns true if the given node is equal (isomorphic) to the current node.
- public void mergeInto(HNode node) – merges the current node into the provided node. The current node should not be used further after this.
- public HNode makePos(Vertex v, Edge elist[]) – creates a new node that is the positive child of the current one.
- public void makeNeg(Vertex v, [Edge elist\[\]](#)) – turns the current node into its own negative child. This is done as an optimization since the node would otherwise be discarded after the negative child was created.
- public boolean isSuccess() – returns true only if the node is successful. A node is considered successful if it represents a state where the source is connected to the target.
- public boolean isFailed() – returns true only if the node is failed. A node is considered failed if it represents a state of the network where the source is unable to connect to the target.
- public FloatingPoint getRel() – returns the reliability of the node
- public int getCode() – returns an integer that represents the state of the node, in case a hash table is to be used

3.6 Qc

This class was implemented as a circular queue in the initial version of the OBDDH, and its name refers to the fact it was the queue storing HNodes at the current level of the OBDDH – the level now being processed. You should implement it in the most efficient way, which is not necessarily a circular queue. The actual order of the nodes in Qc isn't particularly important

Provided: None

Required:

- public Qc(HNode root) – initializes the structure to be empty apart from the given root node.
- public HNode front() – removes the first node in the structure and returns it.
- public boolean isEmpty() – returns true if no nodes remain stored, or false otherwise.
- public void insert(HNode node) – inserts the given node to the back of the structure

3.7 Qn

This class was implemented as a circular queue in the initial version of the OBDDH, and its name refers to the fact that it was the queue storing HNodes at the next level of the OBDDH, one level below the level currently being processed. HNodes are added to Qn one at a time, and only if there isn't an equal node already present. They are removed all at once and moved to Qc by the swap method.

Provided: None

Required:

- public Qn() – initialized the structure to be empty
- public boolean isEmpty() – returns true if no nodes remain stored, or false otherwise
- public int numNodes() – returns the number of nodes currently stored in the structure
- public boolean add(HNode node) – checks to see if a node equal to the provided node is already stored in the structure. If an equal node is found the provided node is merged into the existing one. If no equal node is found the provided node is stored. **It is imperative that this method be as efficient as possible.** This means that deciding whether or not a duplicate node exists must be efficient. This method returns true if an existing node was found or false otherwise.
- public void swap(Qc current) – moves all of the nodes stored in this structure and adds them to *current*, one at a time. This is used when a level of the diagram has been finished and the algorithm is about to start processing the next level. When this method has completed this structure should be empty and all nodes that were in this structure should be in *current*.

3.8 OBDDH

This stores the main diagram and computes the reliability of the network. The main class definition will be provided. It includes a class variable *Network net* that stores the network to be analysed. It also has two private class variables that store HNodes – private Qn next and private Qc current.

Provided:

- public OBDDH(String filename) – default constructor that creates the initial, empty diagram and then loads up the network and computes the network's reliability
- private void compute() – computes the reliability of the network, which must have been initialized before this is called
- public void displayRel() – displays the reliability of the network to the screen

Required:

- `private void loadNetwork(String filename)` - initializes the network by calling the constructor to load information from a file
- `private void addToQn(HNode node)` - after a new child node has been created this method attempts to add it to *next*.

4 Submission

Submit electronically via Blackboard. As per the preamble, make sure to submit early. In fact I strongly encourage you to submit multiple times - I will only mark the last attempt. Take care not to submit your last version late though. Read the submission instructions very carefully.

You should submit a single file, which should be zipped (.zip) or tarred (.tar.gz). Check that you can decompress it on `saeshell01p.curtin.edu.au`. This is also the computer on which your work will be tested, so make sure that your work runs there. The file must be named `DSA_Assignment_<id>` where the `<id>` is replaced by your student id. There should be no spaces in the file name; use underscores as shown.

The file must contain the following:

- Your code. This means all .java files needed to run your program. Do not include any .class files. Do include code provided to you as part of the assignment if that is required to run your program.
- Your test harnesses. One of the easiest ways for us to be sure that your code works is to make sure that you've tested it properly. Our test harnesses may not work for some of your classes and you may have classes that we're not specifically asking for, so make it easy for us to test your work. You do not have to provide a test harness for any code that is provided to you (such as most of HNode and its associated class) but you may want to write one anyway to make sure that you understand how the code is behaving. A test harness for class X should be called `UnitTestX`. For example the test harness for Vertex should be named `UnitTestVertex`.
- Documentation for your code, as described in Section 4.1.
- A signed and dated cover sheet. These are available from the [[Computing Colloquium]] Blackboard unit or from the Computing reception (building 314, 3rd floor). You can sign a hard copy and scan it in or you can fill in a soft copy and digitally sign it.
- You may include a Makefile, but this is not required. We will use `javac *.java` to compile your files and run the unit tests by their expected names.
- Do not include any directories - in other words add files directly to the archive rather than compressing an entire directory. When your submission is expanded, we want to be able to use `javac` in the current directory. We will be using scripts for some of the testing and don't want the scripts failing because your code is not where we expect it to be. If we decompress your submission it should create a number of files in the directory we ran the decompression in - if this isn't the case you could get

zero marks for the script-based marking, so check your submission before uploading it.

- Do not include .class files or anything else that we do not need. We will compile your .java files to ensure that what we're testing is what we're reading.

Make sure that your file contains what is required. Anything not included in your submission will not be marked, even if you attempt to provide it later. It is your responsibility to make sure that your submission is complete and correct.

4.1 Documentation

You need to submit documentation in docx or pdf format. Your documentation should include the following:

- An overview of your code, explaining any questions that the marker may have. This is supplemented by the comments in your code. In general, if the marker is looking at your code and isn't sure why you have done something in that particular way, they will check your documentation to see if they can find an explanation. Using an automated documentation system like Javadocs may be very helpful. It is not required, though.
- A description of any classes you have, other than those covered in the class specification section. If you have additional classes then you need to let us know not only what the purpose of that class is but why you chose to create it. As part of this, also identify and justify any places where it was possibly useful to create a new class but you chose not to, especially when it comes to inheritance.
- Justification of all major decisions made. In particular, when you choose an ADT, underlying data structure or an algorithm, you need to justify why you chose that one and not one of the alternatives. These decisions are going to be of extreme importance in this assignment.

4.2 Marking

Marks will be awarded to your submission as follows:

- [40 marks] Decisions made. This is mainly things like choosing the right class breakdown, ADT, underlying data structure and algorithm. Note that you will only get these marks if you adequately justify your decisions and properly implement them. So, for example, if you choose a sorting algorithm and adequately justify it but aren't actually able to implement it you'll only get part of the marks. Any justification should be presented in the documentation, which means that your documentation is worth 40% of your marks for this assignment. Start documenting early.
- [20 marks] Code testing. We'll have a number of tests to run, and you get marks proportional to how many tests your code passes. If your code passes all of the tests, then you will get all of these 20 marks. This will involve loading a network and computing the reliability. Note that these tests will be script based, so if you have something set up in a way that breaks specifications so that the script doesn't work, you won't get those marks. Some of the tests will be based on the efficiency of your

implementation, so a naïve implementation that doesn't use optimal structures and algorithms may not be able to complete the tests in a reasonable amount of time.

- [15 marks] Implementation of the Network and associated classes. This includes Vertex, Edge and any other classes you create as part of this. It also includes implementing all required methods associated with these classes. We will use unit testing as well as looking at code quality; your test harnesses will have a big impact on these marks.
- [25 marks] Implementation of the OBDDH and associated classes not relating to the Network. It also includes HNode, Qc, Qn and any other classes that you have created to support them, as well as all required methods for these classes. We will use unit testing as well as looking at code quality; your test harnesses will have a big impact on these marks.
- Marks will be deducted for not following specifications outlined in this document, which includes incorrect submission format and content and using built-in Java ADTs.
- Marks will be deducted for changing or not using provided code.
- If the cover sheet isn't provided with your submission, your submission will not be marked and you will be awarded zero (0) marks. If you forget to submit the cover sheet you will be allowed to submit it separately to the unit coordinator (by e-mail or in person) but will lose 5 marks.

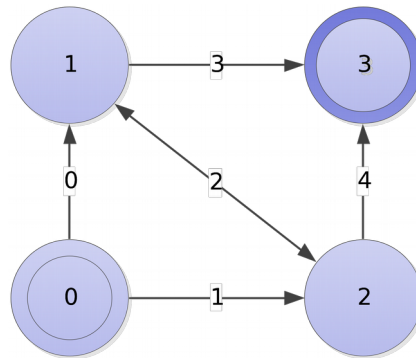
The aims of this marks breakdown is as follows:

- To reward you for good design decisions, but not unless you're able to implement them. This is the core aim of the unit.
- To let you score some marks if you get the program working but make with poor decisions. If all of your decisions are based only on ease of implementation you can still score the 40 marks from the two implementation categories and will score most of the marks from testing (around 15 to 17), meaning you can pass if all of your code works perfectly and is of sufficient quality. It's taking a risk, though.
- To promote good testing. Industry is repeatedly telling us they are really looking for students who can properly test their code.
- To teach you to follow specifications carefully. If you name your files incorrectly you can lose large amounts of marks from our tests failing. You also need to meet the exact requirements for class and method templates, and so on. Follow the specifications in this document carefully.
- To reward students who have been serious about doing their practical sheets.

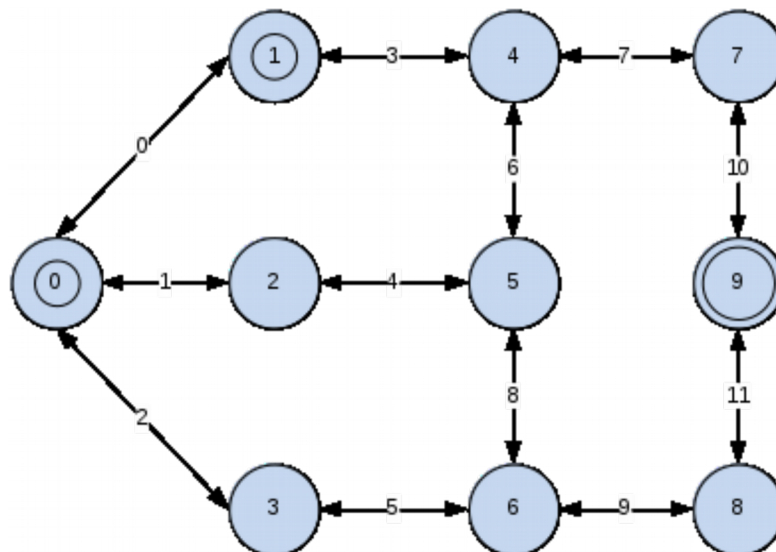
Appendix I: Example of Network Sorting

It is generally believed by reliability researchers that sorting the network in a breadth-first fashion increases the efficiency of the OBDD computation. For this reason, the network read in needs to be sorted appropriately. We sort vertices in increasing order from the source and the edges in increasing order of first the smaller endpoint and then the larger.

Consider the simple network below. The vertices are numbered from 0 to 3 and the edges from 0 to 4; vertex 0 is the source and vertex 3 is the target. A correct ordering of the vertices would be 0,1,2,3 but 0,2,1,3 would also be correct since both vertices 1 and 2 are one edge away from the source and can thus be in any order, but vertex 0 is the source (distance zero) and vertex 3 is two edges away and must thus be last.



The edges for the network are $e_0 = (0,1)$, $e_1 = (0,2)$, $e_2 = (1,2)$, $e_3 = (1,3)$ and $e_4 = (2,3)$. These are in the correct order given the first ordering of the vertices. If we had put vertex 2 ahead of vertex 1 then we would need to switch the order of e_0 and e_1 , as well as the order of e_3 and e_4 .



As an additional example, consider the network above. If the source is vertex 0 (I'll use v_0 to denote this) then the vertices at distance 1 are v_1 , v_2 and v_3 , the vertices at distance 2 are v_4 , v_5 and v_6 , those at distance 3 are v_7 and v_8 and the only vertex at distance 4 is v_9 . The distance in each case is the smallest number of steps (also called hops) from the source to that vertex. You can get to v_4 through v_1 in two hops, so the distance is two despite also being able to reach it through longer paths (for example via v_3 , v_6 and then v_5).

This means that the following are some of the many valid sortings of the vertices:

- v0, v1, v2, v3, v4, v5, v6, v7, v8, v9
- v0, v3, v1, v2, v5, v6, v4, v8, v7, v9

As can be seen, v0 has to be first since it's the source, and thus the only vertex at distance zero. The three distance 1 vertices come next, in any order, followed by the three distance 2 vertices, also in any order. Finally we have the two distance 3 vertices (in any order) and the only distance 4 vertex must be last. Thus there are many possible "correct" sortings.

Once the sorting of vertices is complete, there is generally only one correct sorting of edges. The exception is when there are two edges between the same pair of vertices (for example if the graph above had a second edge between v0 and v1) in which case the order of those edges can vary.

Consider an edge (v_{from}, v_{to}) where v_{from} is the *from* vertex and v_{to} is the *to* vertex. We always write an undirected edge so that the *from* vertex is lower in the sorting that we have ended up with. For example if the edge is from v4 to v5 we would write it as (v4,v5) if using the top sorting shown above and as (v5,v4) for the second sorting. If the edge was directed, we would not have the freedom to change the *from* and *to* vertices.

For a sorting, define the comparison as follows:

- $(a,b) < (c,d)$ if and only if a is before c in the sorting OR $a = c$ and b is before d in the sorting.

For example, with the first sorting we would have $(v0,v1) < (v0,v2)$ and $(v1,v4) < (v3,v6)$. For the second sorting we still have $(v0,v1) < (v0,v2)$ but now $(v3,v6) < (v1,v4)$ since v3 is before v1 in the ordering.

With this inequality, the edges are now sorted in ascending order. The ordering of edges for the two sortings above are:

- e0, e1, e2, e3, e4, e5, e6, e7, e8, e9, e10, e11
- e2, e0, e1, e5, e3, e4, e8, e6, e9, e7, e11, e10

In both cases, there is only one correct ordering. Any other ordering does not follow the guidelines given.

Appendix II: Example of Network File

The network files for the assignment will be provided to you in XML; there are a number of samples available on Blackboard. You will need to read one of these in before the building of the decision diagram can begin. The sample documents show the general format of a network file:

- The network file is enclosed in the `<NET>` and `</NET>` tags. All information relating to the network should be found between these.
- Each vertex is defined in a VERTEX tag that has components NAME, REL (reliability), COST and TYPE. The name and type strings are in double

quotes and the reliability should be interpreted as a floating point number (either float or double, depending on your class choice). The cost is the cost in dollars (US) per unit. You may assume that the vertex name is unique. A vertex can also have either SOURCE or TARGET as an attribute, which defines the communication requirement. We generally have the first vertex as the source vertex and the last be the target, but this isn't always the case and is actually a disadvantage in some cases so don't make that assumption.

- Each edge is defined in an EDGE tag that has the same components as above but also has two additional fields, FROM and TO. Both of these fields are strings that detail the endpoints of the edge. Our edges are assumed to be directed, but an additional -u flag means that the edge is undirected (*i.e.*, can be traversed either way). Each of the FROM and TO strings should match the name of a vertex read in earlier. Note that the cost for edges represents the cost in cents (US) per meter of the component, and this cost can be zero for wireless connections.

The sample files list all vertices first, and then all edges. This is not necessarily required, but a vertex must be described before it can be used as the endpoint for an edge. In general, it's easier to list all vertices first.

Note that this specification applies only to an unsorted network file. You may store sorted network files in any format you wish, but choosing a different format means creating a different load method for sorted files. As with everything, considered relative efficiency.