



Data Structures and Algorithms

Lecture 9: Heaps



Copyright Warning

COMMONWEALTH OF AUSTRALIA

Copyright Regulation 1969

WARNING

This material has been copied and communicated to you by or on behalf of Curtin University of Technology pursuant to Part VB of the Copyright Act 1968 (the Act)

The material in this communication may be subject to copyright under the Act. Any further copying or communication of this material by you may be the subject of copyright protection under the Act.

Do not remove this notice



This Week

- Priority queues
- Heaps
 - Array representation of binary trees
 - Analysis of Heap efficiency
 - MaxHeap vs MinHeap
- HeapSort and comparison to other $O(N \log N)$ sorts

Priority Queues

- We've talked about FIFO queues in earlier lectures
- But there is another type of queue that is also fairly common: the Priority Queue
 - Two operations: **add** and **remove**
 - Items **added** to priority queue with an associated **priority**
 - » Priority indicates how quickly the item must be dealt with
 - » Highest priority item in the queue is always **removed** first
 - Priority-based processing is quite common. Examples:
 - » Task scheduling for CPU execution by an operating system
 - » Inventory ordering: low-stock and/or popular items are the most important (highest priority) to order
 - » Preferential treatment for loyal and/or large customers

Priority Queues – Priority Definition

- The priority value is usually an integer
 - `void add(int priority, Object value)`
 - » Could be a float, but that's less common
- But what constitutes “high priority”? Two options:
 - Higher integer values = higher priority
 - » *e.g.*, bigger vs smaller
 - Lower integer values = higher priority
 - » *e.g.*, first, second, third: like a race
 - These lectures will assume high value = high priority
 - » Just makes it easier to keep it straight in your head!

Priority Queues – Implementation

- So how can we implement a priority queue ADT?
- So far, we only know of arrays and linked lists. Both have a fairly similar priority queue implementation:
 - **Add**: add them in sorted order according to priority
 - » Requires searching through array/list to find insertion point
 - » Averages $N/2$ steps, ie: **Add** = $O(N)$
 - **Remove**: simply take from the rear (since highest-priority will be at the end when in sorted order)
 - » Fast: **Remove** = $O(1)$
 - » Note: We never remove anything but the highest-priority item

Priority Queues – Implementation

- An alternative is to avoid sorting the data and instead make it remove's problem to find the highest priority
 - **Add**: Append the item to end of array/list
 - » Fast: **Add** = $O(1)$
 - **Remove**: Search through list to find highest-priority item
 - » Must go through all N items just in case highest is last item
 - » *i.e.*, **Remove** = $O(N)$
- Whichever alternative is taken, you cannot avoid having one of add or remove being $O(N)$
 - Can we do better than $O(N)$? Fortunately, yes: that's what a Heap data structure is for

Heaps

- The **heap data structure** is *not* the same as “the heap” used in programming languages to denote the area of memory used to allocate objects
- Heaps are organised in a binary tree (but NOT as a binary search tree) where the highest priority item is at the root, and lower-priority items are below
 - Requirement: children are always smaller than their parent
 - » *i.e.*, a heap organises items (weakly sorted) from top to bottom
 - Thus it is **NOT** organised like a binary **search** tree, which requires that $\text{leftChild} < \text{parent} < \text{rightChild}$
 - » *i.e.*, a binary search tree organises items from left to right


Heaps and Priority

- Since heaps are so closely associated with priority queues, they also explicitly define priority order
 - A **max heap** is a heap where a **larger** priority value is considered a higher priority
 - A **min heap** is a heap where a **smaller** priority value is considered a higher priority
- For the remainder of this lecture we will be working with max heaps (unless otherwise stated)

Heap Binary Tree – Properties

- The main constraint that a heap tree has is that each child must be of lower priority than its parent
 - This guarantees that the highest priority item is the root
 - It doesn't matter if the left child is larger or equal to the right child, or vice-versa
- By a little bit of clever algorithm design, the heap is also guaranteed to be always **almost-complete**
 - Thus always guaranteeing $O(\log N)$ access time
 - We will see how this is guaranteed when we discuss how `add()` and `remove()` work in a heap

(Max) Heap – Example



Heaps – Some Notes

- A heap mandates that children nodes are always of lower priority than their parents
 - This is enough to guarantee that the root is the highest priority item, which is enough for a priority queue
 - Ordering is vertical, but higher-priority items only *tend* to be higher up in the tree
 - » Different subtrees may contain much different priorities
 - » *e.g.*, 6 is lower in the tree than 5, but is of higher priority
 - Thus a heap is only *weakly* ordered
- Heaps can also contain duplicate priority items
 - Priority is *not* a unique key for lookup!

Array Representation of Binary Trees


- Let's take a small detour and discuss different ways of representing a binary tree
- Normally, trees are represented (implemented in memory) via tree nodes pointing at other tree nodes
 - Nodes are scattered about in memory (ie: non-contiguous)
 - Each node has left child and right child pointers
- But it is also possible to represent a binary tree with an array

Array Representation of Binary Trees

- There are a few ways to go about representing a tree in an array form
- Heaps use a form that has certain desirable properties
 - » Other forms just complicate things
 - Heaps consider the tree as a set of levels, and ‘pack’ the levels into an array, one level after the other
 - This works *only* because it is almost-complete
- Converting a heap’s binary tree to array form is easy:
 - Simply read off the tree level-by-level and build the array in that order

Heap Array – Example

0	9
1	8
2	5
3	6
4	0
5	2
6	3



Heap Arrays

- This array form has a crucial benefit: it allows us to *calculate* how to go up and down the tree via arithmetic
 - The root is at element [0] in the array
 - All siblings are beside each other in the array
 - Thus if we are at node [currIdx], then:
$$\text{leftChildIdx} = (\text{currIdx} * 2) + 1$$
$$\text{rightChildIdx} = (\text{currIdx} * 2) + 2$$
$$\text{parentIdx} = (\text{currIdx} - 1) / 2$$
 - The $* 2$ comes about since we have a *binary* tree
 - parentIdx is derived by inverting equation for leftChildIdx
 - » Inversion of rightChildIdx is equivalent since $/ 2$ is DIV 2 and the right child index is always an even number


Heap Arrays

- Why does it matter to use arithmetic for traversal?
 - Because as we will see later, a heap needs to be able to traverse up *and* down the tree
 - In a tree form, this would require the addition of a ‘parent’ pointer in each node – extra memory overhead
 - With the arithmetic-based traversal, we can even do away with the left/right child pointers: **no memory overhead!**
 - » ie: we only need to store the priority+data in the array
- BUT: the arithmetic only works for [almost-]complete trees
 - All levels are full (*i.e.*, exactly 2x larger than parent level), except for the last level which is filled from the left

Heap – Add

- Strategy: Initially place a new item in the next slot of the almost-complete tree
 - This guarantees the tree will remain almost-complete
 - The ‘next slot’ is easy to find: it’s at the end of the used portion of the array!
- Then ‘trickle’ the new item up through the tree until it meets a parent of equal or higher priority
 - Trickle-up = swapping based on priority checks vs parent
 - Essentially, we promote the new item until it reaches the place where it should be at (according to priority)
 - It doesn’t matter what branch it starts in: remember, heaps are only weakly ordered

Add Example



Details of Add's Trickle-Up

- `add()` is essentially a loop that swaps the new node up the tree (trickle-up) while the following conditions hold true:
 - The new node has NOT made it to the root, AND
 - The parent's priority is lower than the new node
- Trickle-up can be done iteratively or recursively.

Iterative Trickle-Up

```
IMPORT heapArray, curIdx
```

```
EXPORT heapArray
```

```
Assertion: WHILE cur NOT root AND cur > parent DO
```

```
    Swap cur with parent, then try again
```

```
    parentIdx = (curIdx-1)/2
```

```
    WHILE curIdx > 0 AND heapArr[curIdx] > heapArr[parentIdx]
```

```
        temp = heapArr[parentIdx]
```

```
        heapArr[parentIdx] = heapArr[curIdx]
```

```
        heapArr[curIdx] = temp
```

```
        curIdx = parentIdx
```

```
        parentIdx = (curIdx-1)/2
```

```
    ENDWHILE
```

Recursive Trickle-Up

```
IMPORT heapArray, curIdx
```

```
EXPORT heapArray
```

```
Assertion: IF cur NOT root AND cur > parent THEN
```

```
    Swap cur with parent, then try again
```

```
parentIdx = (curIdx-1)/2
```

```
IF curIdx > 0 THEN
```

```
    IF heapArr[curIdx] > heapArr[parentIdx] THEN
```

```
        temp = heapArr[parentIdx]
```

```
        heapArr[parentIdx] = heapArr[curIdx]
```

```
        heapArr[curIdx] = temp
```

```
        trickleUp <- heapArray, parentIdx
```






Heap – Remove

- Since the heap is a way of implementing a priority queue, ie: we always want to remove the item with the highest priority
 - And the heap is organised such that the highest priority item is *always* the root node
- It then follows that we will always remove the root node
 - Of course now we have a problem – we have lost our root node and our tree is not almost-complete anymore

Heap – Remove

- Strategy: Take a copy of the root at element [0], and move the last element to replace the root
 - Since the last element is at the final almost-complete position, removing it will maintain almost-complete tree
 - But now the root is going to be a low-priority item
 - » *i.e.*, the heap's rule that $\text{parent} \geq \text{children}$ is being violated
- So 'trickle' this incorrect root node down through the tree until it finds its correct position
 - *i.e.*, swap down until *neither* child is higher priority
 - This often involves swapping to the bottom of the tree

Remove Example



Details of Remove's Trickle-Down

- Removing the root and moving the last node into the root's position is pretty easy
 - Just copy the root to a temp variable, and copy the last-used element in the array to the root at [0]
- After that, trickle-down is quite similar to trickle-up: keep trickling-down the node while:
 - The node still has children (ie: $\text{currIdx} < \text{count}/2$) AND
 - *Either* children's priority is higher than the node

Details of Remove's Trickle-Down

- However, unlike `add()`, `remove()` has two possibilities for swapping:
 - Swap with left child OR Swap with right child
- We must swap with the higher-priority child to maintain that “all parents are higher priority than children”
 - To simplify the code: before the swap, compare the two children *first* and choose the highest-priority child
 - *Then* compare the trickling node with that child and swap if the child that has higher priority

Iterative Trickle-Down

```
IMPORT heapArray, curIdx, numItems
EXPORT heapArray

lChildIdx = curIdx * 2 + 1
rChildIdx = lChildIdx + 1
keepGoing = true
WHILE keepGoing AND lChildIdx < numItems      //is a left child
    keepGoing = false
    largeIdx = lChildIdx
    IF rChildIdx < numItems                    //is a right child
        IF heapArr[lChildIdx] < heapArr[rChildIdx]
            largeIdx = rChildIdx              //find largest child
    IF heapArr[largeIdx] > heapArr[curIdx]
        swap <- heapArr, largeIdx, curIdx
        keepGoing = true
    curIdx = largeIdx
    lChildIdx = curIdx * 2 + 1
    rChildIdx = lChildIdx + 1
ENDWHILE
```

Recursive Trickle-Down

```
IMPORT heapArray, curIdx, numItems
```

```
EXPORT heapArray
```

```
lChildIdx = curIdx * 2 + 1
```

```
rChildIdx = lChildIdx + 1
```

```
IF lChildIdx < numItems                                     //is a left child
```

```
  largeIdx = lChildIdx
```

```
  IF rChildIdx < numItems                                   //is a right child
```

```
    IF heapArr[lChildIdx] < heapArr[rChildIdx]
```

```
      largeIdx = rChildIdx                                 //find largest child
```

```
IF heapArr[largeIdx] > heapArr[curIdx]
```

```
  swap <- heapArr, largeIdx, curIdx
```

```
  trickleDown <- heapArray, largeIdx, numItems
```

Heaps – Complexity Analysis

- Add: Best = $O(1)$, Average/Worst = $O(\log N)$
 - Best case: occurs when adding a very low priority item
 - » It won't be trickled up since it is already in the right spot
 - » Thus $O(1)$
 - Worst case: occurs when the added item has highest priority and must be trickled up all the way to the root
 - » Since the heap tree is *always* almost-complete, there are $\log N$ levels to trickle through, resulting in $O(\log N)$
 - Average case: $\frac{1}{2}$ the items are at the bottom of the tree.
 - » Which is $O(\log N)$

Heaps – Complexity Analysis

- Remove: $O(\log N)$ for all cases
 - Because we take the last node (which will be among the lowest priorities), place it at the root and trickle down
 - » Even in the best case there must be some trickle-down since the node's correct place was at the very bottom of the tree
 - » And in fact it will usually trickle *all* the way down!
 - Thus $O(\log N)$ for pretty much all cases, even best case

Heaps – Summary

- Data is stored in a weakly-ordered way
 - There is some order (parent larger than children), but nothing like in a BST, thus ordered traversal of the tree is not possible
- Only the first item (root) can be taken
 - Heaps are not useful for searching for a particular value
 - » We aren't storing by key, we are storing by *priority*
- Stores the binary tree in an array form and uses arithmetic on element indexes to traverse the tree
 - And the tree is always in an [almost]-complete state
- *Both* add and remove are fast $O(\log N)$ operations
 - Plus add/remove are crafted to maintain almost-completeness



HeapSort

- A heap returning items in priority order is kind of like getting data in sorted order, just one at a time
- This implies we can use heaps to perform sorting
 - Take an array of unsorted data
 - Add all elements of the unsorted array into a heap, using the element value as the priority
 - » This will organise the elements into a heap
 - Remove each element from the heap one at a time and place them back into the array
 - » Since a heap returns highest-priority first, the elements will come out in sorted order (or reverse sorted order)

HeapSort

- Depending on whether you are using a max-heap or a min-heap will affect the order of the sort
 - Max-heaps will return larger values first, hence the heap is effectively providing data in reverse order
 - » Not a big deal: simply populate the target array in *reverse* order (from back to front).
 - This is just as efficient as using a min-heap and populating the target array in forwards order
 - » The only difference between the two is that you either loop from $0 \dots N$ (min-heap) or loop from $N \dots 0$ (max-heap)

HeapSort Time Complexity

- If a heap is available HeapSort is a particularly simple algorithm to implement:
 - An initial for loop to add all array values to the heap
 - A second for loop to take them all out one at a time
- But how efficient is it?
 - **Add:** $O(\log N)$ done N times = $O(N \log N)$
 - » OK, best case of $O(1) * N = O(N)$, but that's rare!
 - **Remove:** $O(\log N)$ done N times = $O(N \log N)$
 - Total = Add + Remove
 - = $O(N \log N) + O(N \log N)$ (or best case $O(N) + O(N \log N)$)
 - = $O(N \log N)$

In-Place HeapSort

- Hence HeapSort is as scalable as Quick+MergeSort
 - Unfortunately, it is **unstable** since we may get equal-priority values being swapped relative to each other
 - The simple approach outlined is also **not in-place**
 - » The heap has an array that is the same size as the original array
 - However, it is possible to make HeapSort in-place by integrating it into the heap's code (**more complicated!**)
 - » First organise the array into a heap incrementally by 'expanding' the heap one element at a time and adding that new element
 - This is termed to 'heapify' the array
 - » Then every time the root is taken from the heap, add it to the array slot that has just been 'vacated' by the last node

heapify

```
IMPORT heapArray, numItems
```

```
EXPORT heapArray
```

```
ASSERTION: imported array will be random, exported will  
           be a heap
```

```
// start at last non-leaf, go backwards
```

```
  for ii = (numItems/2)-1 downto 0          //0 based array
```

```
    // put iith element in correct place in heap
```

```
    trickleDown <- heapArray, ii, numItems
```

heapSort (in-place)

```
IMPORT array, numItems
```

```
EXPORT sortedArray
```

```
ASSERTION: imported array will be random, exported will  
           be the same array sorted
```

```
heapify <- array, numItems
```

```
for ii = numItems-1 downto 1           //0th item will be sorted
```

```
    swap <- array, 0, ii
```

```
    trickleDown <- heapArray, 0, (ii)   //ii is numItems--
```

heapSort example

Import

0	1	2	3	4	5	6	7	8
12	16	3	11	10	1	2	5	4

**After
Heapify**

0	1	2	3	4	5	6	7	8
16	12	3	11	10	1	2	5	4

**After 1st
swap**

0	1	2	3	4	5	6	7	8
4	12	3	11	10	1	2	5	16



After
trickleDown

0	1	2	3	4	5	6	7	8
12	11	3	5	10	1	2	4	16

After 2nd
swap

0	1	2	3	4	5	6	7	8
4	11	3	5	10	1	2	12	16

After
trickleDown


0	1	2	3	4	5	6	7	8
11	10	3	5	4	1	2	12	16

After 3rd
swap

0	1	2	3	4	5	6	7	8
2	10	3	5	4	1	11	12	16

After
trickleDown

0	1	2	3	4	5	6	7	8
10	5	3	2	4	1	11	12	16



After 4th
swap

0	1	2	3	4	5	6	7	8
1	5	3	2	4	10	11	12	16

After
trickleDown


0	1	2	3	4	5	6	7	8
5	4	3	2	1	10	11	12	16

After 5th
swap

0	1	2	3	4	5	6	7	8
1	4	3	2	5	10	11	12	16

After
trickleDown

0	1	2	3	4	5	6	7	8
4	2	3	1	5	10	11	12	16



After 6th
swap

0	1	2	3	4	5	6	7	8
1	2	3	4	5	10	11	12	16

After
trickleDown

0	1	2	3	4	5	6	7	8
3	2	1	4	5	10	11	12	16

After 7th
swap

0	1	2	3	4	5	6	7	8
1	2	3	4	5	10	11	12	16

After
trickleDown

0	1	2	3	4	5	6	7	8
2	1	3	4	5	10	11	12	16

After 8th
swap

0	1	2	3	4	5	6	7	8
1	2	3	4	5	10	11	12	16

HeapSort

- ✓ Can be an in-place sort if built into the Heap class
- ✓ Consistently $O(N \log N)$ for all cases
- ✓ By far the easiest $O(N \log N)$ algorithm to implement *if* you can make use of an *existing* Heap class
 - Just a couple of for loops: one to insert all the elements, another to extract them out in [reverse-]sorted order
 - Although with this approach it cannot be made in-place
- ✗ Unstable sort
- ✗ Poor use of a modern CPU's L2 cache
 - Trickle-up and trickle-down jump all over the array
- ✗ Requires implementing a Heap:

MergeSort

- ✓ Easy to make execute in parallel
 - Since different split and merge branches are independent, we can assign each branch to a different CPU
- ✓ Makes efficient use of a modern CPU's L2 cache
 - It merges two sub-arrays that are *beside* each other
 - » AND goes through each array from left to right
 - So accesses are always close together: perfect for L2 caching
 - » With a large L2 cache, MergeSort can become very fast:
L2 accesses are up to 5x faster than main memory accesses
- ✓ And: Stable, consistently $O(N \log N)$ for all cases
- ✗ But: Not an in-place sort

QuickSort

- ✓ Easy to make execute in parallel
 - Since different split+partition branches are independent, we can assign each branch to a different CPU
- ✓ Makes some use of a modern CPU's L2 cache
 - Splitting will mean that it eventually operates on sub-arrays that are small enough to fit into the L2 cache
 - But not as good as MergeSort at this
- ✓ **And:** in-place sort
- ✗ **But:** Unstable sort, recursive (stack overflows), $O(N^2)$ worst-case, fairly complicated to implement well

Next Week

– Advanced Trees (Red-Black, 2-3-4, B-Trees)

